

Veltrixa Cloud

Data Flow Mapping

AI Document Classification Service

FICTIONAL PORTFOLIO CASE STUDY
This document demonstrates a simulated privacy-governance workflow. It is not an operational approval record.

Document ID	VG-PGO-DFM-001
Version	1.2
Classification	Portfolio Sample
Prepared by	Privacy Governance Analyst
Document owner	Privacy Governance Office (fictional)
Date	August 2026

Document Control

0.1 Document Information

Field	Value
Document Title	Data Flow Mapping – AI Document Classification
Document ID	VG-PGO-DFM-001
Status	Portfolio Sample – Final Draft
Business Owner	Product Management (fictional)
Technical Owner	Cloud Engineering (fictional)
Review Frequency	Annual or following a material architectural change

0.2 Version History

Version	Date	Author	Description
0.1	July 2026	Privacy Governance Analyst	Initial portfolio document
1.0	July 2026	Privacy Governance Analyst	First complete portfolio edition
1.1	August 2026	Privacy Governance Analyst	Architecture, transfer, retention and traceability revision
1.2	August 2026	Privacy Governance Analyst	Data-flow, regionality, backup, deletion, inventory and validation corrections

0.3 Review Workflow

Role	Portfolio Status
Lead Solution Architect	Illustrative review role – no actual approval
Senior Privacy Counsel	Illustrative review role – no actual approval
Head of Privacy	Illustrative review role – no actual approval

0.4 Related Documents

- Privacy Discovery Questionnaire
- Privacy by Design Review
- DPIA Screening
- ROPA
- Vendor Reviews
- Transfer Impact Assessment
- Retention Schedule
- Privileged Support Access Procedure

1. Introduction

1.1 Executive Summary

Veltrixa Cloud is a fictional enterprise SaaS platform providing AI-assisted document management. This Data Flow Mapping records how personal data enters, moves through, is transformed within, is accessed in and is deleted from the AI Document Classification capability. It provides a factual baseline for Privacy by Design reviews, DPIA screening, ROPA maintenance, vendor assessment and transfer analysis.

1.2 Purpose

The document establishes a shared, testable understanding of the personal-data lifecycle and identifies assumptions that engineering, security, privacy, product and vendor-management stakeholders would validate in a real implementation.

1.3 Scope

In Scope

Authentication, document upload, validation, malware handling, temporary processing, AI classification, metadata generation, human review, permanent storage, search, access, logging, vendor telemetry, backup, retention and deletion.

Out of Scope

Customer business processes outside Veltrixa Cloud, customer legal bases for the underlying document content, model development and training, and detailed incident-response procedures.

Assumptions and constraints

- Customer document content and AI inference workloads are designed to remain in AWS eu-central-1 (Frankfurt).
- The regional scope of identity data, operational telemetry, vendor support access and service metadata is validated separately for each service and provider.
- Customers may upload documents containing special-category data or Article 10 data. The platform therefore combines customer instructions and guidance with sensitivity-agnostic technical safeguards.
- The AI service does not use customer documents for model training and does not transmit document content to external AI providers.

1.4 Applicable Regulations and Standards

Framework	Relevance
GDPR	Articles 5, 9, 25, 28, 30, 32, 35 and 44–49
EU AI Act	Applicability and risk-classification screening; the use case is not assumed to be high-risk
ISO/IEC 27001:2022	ISMS risk management and information-security controls
ISO/IEC 27701:2025	Privacy information management and accountability
NIST Privacy Framework 1.0	Current final framework; Privacy Framework 1.1 Initial Public Draft monitored

2. Business Context

2.1 Business Process

Authorised customer users upload business documents through a secure web application. The platform validates and scans each file, stores a temporary encrypted copy, analyses the content using an internally hosted AI Classification Service and produces draft document types, tags and confidence scores. Users review and confirm or modify the suggestions before documents and approved metadata enter permanent storage and searchable workflows.

2.2 Data Subjects

Data Subject	Typical Examples
Employees	Employment contracts, HR records
Job applicants	CVs and cover letters
Customer contacts	Business correspondence and agreements
Supplier contacts	Contracts and invoices
Consultants	Statements of work
Other individuals	References, emergency contacts and meeting participants
Children or vulnerable persons	Benefits, education or family-related documents, where applicable
Customer users and administrators	Named user accounts, authentication events, IP addresses, search activity and audit records

2.3 GDPR Roles

Veltrixa acts as processor for customer document content and related metadata processed under the customer's documented instructions. Veltrixa acts as controller only for processing activities for which it independently determines the purposes and essential means, such as defined account-administration, security-monitoring, abuse-prevention and service-management activities. Role allocation is assessed per processing activity rather than assigned once for the platform as a whole.

AWS, Datadog and the external Identity Provider are classified by processing context. They may act as subprocessors for customer-directed processing and as processors for operational personal data controlled by Veltrixa. Vendor roles are assigned per processing activity. Any provider processing for its own billing, account-management, security or legal purposes is recorded separately in the vendor review and is not treated as customer-directed subprocessing without evidence.

2.4 Categories of Personal Data

Category	Examples
Identity and contact	Name, address, email address, telephone number
Employment	Job title, employer, employee ID
Financial	Invoice details and payment references
Technical and usage	User ID, IP address, session identifiers, timestamps and audit events
Special categories – possible	Health, trade-union membership, political opinions, religion, ethnicity, biometric or genetic data
Article 10 data – possible	Criminal-conviction or offence information contained in uploaded documents
System-generated	Draft and confirmed classifications, tags, confidence scores and workflow metadata relating to an identifiable person

3. Solution Overview

3.1 Architecture Overview

The target architecture separates temporary processing, approved document storage, metadata, identity services and operational telemetry. Customer document content, temporary copies, extracted text and AI inference workloads are designed to remain within regional AWS services in eu-central-1. AWS IAM and any other global or control-plane services are assessed separately and are not represented as region-bound. Customer authentication is provided through an EU-configured external Identity Provider. For this target design, an allowlisted telemetry subset is forwarded from CloudWatch to the Datadog EU site. Document bodies, extracted text and sensitive search terms are excluded before forwarding.

Veltrixa Cloud - Target Architecture

Regional AWS data plane in eu-central-1; global/control-plane services assessed separately




Figure 1. High-Level Architecture. Target components, data stores and supporting services.

3.2 Core Components

Component	Purpose	Personal Data
Web Portal	Upload, review, search and retrieval	Yes
External Identity Provider	Customer authentication and identity events	Yes
Upload Service	Validation, malware scanning and routing	Yes
Temporary S3 Bucket	Short-lived encrypted processing copy	Yes
AI Classification Service	Internal inference and draft metadata generation	Yes
Amazon S3	Approved document storage	Yes
Amazon RDS for PostgreSQL	Metadata and search indexes	Yes
AWS IAM	Workload authentication and authorisation	Limited identifiers
Amazon CloudWatch	Logs, metrics and alerts	Limited technical data
Datadog EU	Minimised operational telemetry	Limited technical data
AWS Backup	Recovery copies of approved S3/RDS data	Yes

4. Data Flow Mapping

Figure 2 and the following stage tables describe the normal end-to-end lifecycle. Supporting and exceptional flows are recorded where they materially change access, location, retention or recipients.

Personal Data Flow - Nine Processing Stages

The normal lifecycle is shown below; exception and supporting flows are identified separately.




Figure 2. Personal Data Flow. *Nine-stage lifecycle from upload to verified deletion.*

4.1 Stage 1 – Secure Document Upload

Business Purpose	Receive documents from authenticated customer users.
Description	The user authenticates through the external Identity Provider and uploads a supported document through the Web Portal over HTTPS/TLS.
Personal Data	Document content, user ID, tenant ID, IP address, timestamp and file metadata.
Data Location	Browser-to-platform transmission; temporary receipt by the Upload Service in eu-central-1.
Actors / Access	Authenticated customer user, Identity Provider and Upload Service.
Security Controls	TLS; federated authentication through the external IdP; MFA enforced according to the applicable tenant authentication policy; tenant context; file-size limits; and request validation.
Privacy Observation	The platform does not request document content beyond the file and metadata required for the customer-directed service. Customers remain responsible for the appropriateness of uploaded content, supported by contractual terms and user guidance.

4.2 Stage 2 – Validation, Malware Scanning and Quarantine

Business Purpose	Confirm that uploaded files are supported and safe to process.
Description	The Upload Service validates file type and size and performs internal malware scanning. Rejected or suspected files are isolated from the normal workflow.
Personal Data	Original document, file metadata, malware result and file hash.
Data Location	Encrypted processing or quarantine area in eu-central-1.
Actors / Access	Upload Service and restricted security personnel for approved investigations.
Security Controls	Validation rules, malware scanning, isolated quarantine permissions and audit logging.
Exception / Retention	Unsupported files are deleted immediately after rejection. Files flagged as suspected malware are isolated in quarantine and deleted within 24 hours unless an approved security-incident hold applies. Any hold must record its owner, reason, maximum duration and release trigger. A content-free security event may be retained for 90 days.

4.3 Stage 3 – Temporary Processing

Business Purpose	Prepare the validated document for AI analysis.
Description	A short-lived encrypted copy is stored while classification is pending.
Personal Data	Original document and processing identifiers.
Data Location	Temporary S3 processing bucket in eu-central-1.
Actors / Access	Upload Service and AI Classification Service roles only.
Security Controls	SSE-KMS, bucket policy, service-role isolation, no backup and access logging.
Deletion	Application deletion follows successful or cancelled processing; a lifecycle rule provides fallback expiry within 24 hours for failed or abandoned jobs.

4.4 Stage 4 – AI Classification

Business Purpose	Identify document type and propose metadata.
Description	The internal AI service analyses document structure and text. Extracted text exists only in service memory or ephemeral processing space.
Personal Data	Document content, extracted text and processing identifiers.
Data Location	Internal AI workload in eu-central-1.
Actors / Access	Automated service role; no routine human access.
Security Controls	Internal deployment, service authentication, ephemeral workspace, no external AI provider and no customer-content training.
Privacy Observation	The service produces advisory output and does not itself make a decision producing legal or similarly significant effects.

4.5 Stage 5 – Metadata Generation

Business Purpose	Generate searchable draft metadata.
Description	The AI service returns a proposed document type, tags and confidence score to the application.
Personal Data	Draft classification, tags, confidence score and workflow metadata.
Data Location	Application service and draft records in Amazon RDS.
Actors / Access	Application service and authorised customer users.
Security Controls	Tenant-scoped RBAC, audit logging and separation from document storage.
Privacy Observation	Draft output is treated as personal data where it relates to an identifiable person and remains distinguishable from user-confirmed metadata.

4.6 Stage 6 – User Review and Approval

Business Purpose	Provide meaningful human oversight before metadata is committed.
Description	An authorised user reviews, confirms or modifies the AI-generated suggestions.
Personal Data	Document preview, draft output, confirmed metadata and reviewer audit event.
Data Location	Web Portal and Amazon RDS.
Actors / Access	Tenant-authorised customer users.
Security Controls	Authentication, tenant-scoped RBAC, audit trail and visible confidence information.
Privacy Observation	Whether Article 22 GDPR is engaged depends on the downstream use and effects of the confirmed output; the product does not assume that any human click automatically constitutes meaningful review.

4.7 Stage 7 – Permanent Storage

Business Purpose	Store approved business records and searchable metadata.
Description	After approval, the original document is committed to Amazon S3 and confirmed metadata and search indexes are stored in Amazon RDS.
Personal Data	Document, confirmed metadata, search index entries and audit event.
Data Location	Amazon S3 and Amazon RDS in eu-central-1.
Actors / Access	Tenant-authorised users and restricted platform roles.
Security Controls	SSE-KMS, database encryption, tenant-scoped authorisation, audit logging and backups.
Privacy Observation	Document content and metadata are separated, but both remain personal data where they relate to identifiable individuals.

4.8 Stage 8 – Access and Retrieval

Business Purpose	Allow authorised users to search for and retrieve approved records.
Description	The application queries tenant-scoped metadata in Amazon RDS and retrieves the corresponding document from Amazon S3.
Personal Data	Search terms, metadata, document content, user ID and access event.
Data Location	Web Portal, Amazon RDS, Amazon S3 and CloudWatch audit logs.
Actors / Access	Customer users, customer administrators and exceptionally approved support engineers.
Security Controls	Tenant isolation, object-level authorisation, least privilege, time-bound support access and access logging.
Privacy Observation	Search terms and access histories may themselves reveal sensitive information and are subject to minimisation and retention controls.

4.9 Stage 9 – Retention and Deletion Assurance

Business Purpose	Apply customer instructions and purpose-specific retention schedules.
Description	Customer-configured rules trigger deletion from active document storage, metadata and search indexes. Non-current object versions, backups, logs and failed-processing copies follow separate documented schedules. A restore may be initiated only by a restricted recovery role under an approved recovery event. Before tenant access is reopened, restored records are reconciled against tenant status, retention rules and deletion records, and any deletion instructions effective after the recovery point are re-applied and logged.
Personal Data	Documents, metadata, derived indexes, versions, backups and related logs.
Data Location	Amazon S3, Amazon RDS, AWS Backup, CloudWatch and Datadog EU.
Actors / Access	Automated lifecycle services and authorised administrators.
Security Controls	Deletion workflows, S3 lifecycle configuration, backup expiry, log retention, exception monitoring and deletion evidence.
Privacy Observation	Deletion from the active service may precede expiry from recovery copies. The applicable delay is disclosed and validated against the relevant retention schedule.

Supporting flow – Backup restore and reconciliation: A restore may be initiated only by a restricted recovery role under an approved recovery event. Before tenant access is reopened, the service reconciles restored records against tenant status, retention rules and deletion records, then re-applies any deletion instructions that became effective after the recovery point was created. The restore and reconciliation results are logged.

5. Data Inventory

The inventory distinguishes source data, transient transformations, system-generated metadata and operational records. Retention values are target-design assumptions for this fictional case and would require validation against product configuration and contractual commitments.

Data Element	Source	Purpose	Storage	Primary Access	Retention
Original document	Customer upload	Classification and storage	Temporary S3; approved copy in S3	Tenant users; services	Customer policy; active deletion workflow
Extracted text	Original document	AI inference	Memory/ephemeral workspace	AI service role	Processing session only

Data Element	Source	Purpose	Storage	Primary Access	Retention
Draft classification and confidence score	AI service	Human review	Amazon RDS	Application; reviewer	Until approval, rejection or workflow expiry
Confirmed metadata and search index	User-approved output	Search and workflow	Amazon RDS	Tenant application/RBAC	Customer policy
Upload and processing metadata	Platform	Routing and troubleshooting	Amazon RDS/CloudWatch	Platform and support roles	90 days unless incident hold applies
User account and identity events	User/IdP	Authentication	External IdP; tenant account record in Amazon RDS	User and administrators	Account term plus 30 days
Audit trail	Platform	Security and accountability	CloudWatch	Security and authorised administrators	365 days
Minimised telemetry	CloudWatch	Monitoring and alerting	Datadog EU	Security/operations	90 days
Quarantine copy	Upload Service	Security investigation	Isolated temporary storage	Restricted security role	Delete when no longer required; fallback ≤24 hours
Backup copies	S3/RDS	Recovery	AWS Backup vault	Restricted recovery roles	35 days
Search query and retrieval event	Customer user / platform	Tenant search, retrieval and security audit	Amazon RDS and/or CloudWatch as validated	Tenant application; authorised security/support roles	Search query: 90 days or shorter; access audit event: 365 days

6. Data Stores, Retention and Deletion

Repository	Content	Region	Protection	Backup	Deletion / Expiry
Temporary S3	Processing and quarantine copies	Frankfurt	SSE-KMS; isolated roles	No	Unsupported files deleted immediately; suspected-malware quarantine ≤24 hours unless an approved, recorded incident hold applies

Repository	Content	Region	Protection	Backup	Deletion / Expiry
Amazon S3	Approved documents and object versions	Frankfurt	SSE-KMS; tenant-scoped access	AWS Backup recovery points — 35 days	Active object unavailable when the deletion workflow completes; noncurrent versions permanently removed within 7 days; expired delete markers removed by lifecycle policy; incomplete multipart uploads aborted separately; AWS Backup recovery points expire under the validated backup plan
Amazon RDS for PostgreSQL	Draft/confirmed metadata and search indexes	Frankfurt	KMS encryption; RBAC	RDS automated backups/ PITR — validated configured period; AWS Backup recovery points — 35 days	Application rows and indexes deleted by workflow; native automated backups expire under the RDS retention setting; AWS Backup recovery points expire under the backup plan; manual or copied snapshots prohibited or separately inventoried with enforced expiry
Amazon CloudWatch	Audit logs, security events, metrics	Frankfurt	Encryption; restricted access	No separate backup	90 days telemetry; 365 days audit events
Datadog EU	Minimised operational telemetry	EU site	TLS; access controls; content exclusion	Policy-based	90 days
External Identity Provider	Account and authentication data	EU configuration	MFA/federation; provider controls	Provider-specific	Account term plus 30 days; provider schedule validated
AWS Backup Vault	S3/RDS recovery points	Frankfurt	Encrypted vault; restricted recovery role	Not applicable	Automatic expiry after 35 days; deletion exceptions monitored

Deletion assurance note: S3 versioning, lifecycle rules, incomplete multipart upload cleanup, RDS native recovery periods, AWS Backup recovery-point expiry and any manual or copied snapshots must be validated separately and together. A successful deletion request makes the active record unavailable through the tenant service but is not represented as immediate physical erasure from every version or recovery copy.

Data-subject-rights dependency: For customer-controlled document content, Veltrixa supports the customer's documented instructions by locating records through tenant ID, document ID, object version and linked metadata identifiers. Access, correction, export or deletion may require coordinated action across Amazon S3, Amazon RDS, search indexes and relevant audit records. Backup copies are not restored solely to answer a request; recovery copies remain inaccessible until scheduled expiry unless a documented legal or security exception applies.

7. Data Movement Summary

7.1 Internal Flows

Document content moves between the Web Portal, Upload Service, temporary S3 bucket, internal AI service, Amazon S3 and Amazon RDS over authenticated and encrypted channels. AWS IAM authorises workload access. CloudWatch receives minimised platform events and must not receive document bodies or extracted text.

7.2 External Flows

Flow	Data	Recipient	Safeguard	Validation
Authentication	Account identifiers and authentication events	External IdP – EU	DPA, MFA/federation, EU configuration	Provider region and support access
Monitoring	Minimised telemetry	Datadog EU	DPA, EU site, field filtering and RBAC	Support/onward access and retention
Cloud infrastructure	Documents, metadata and service data	AWS	DPA, encryption, regional service selection	Service-specific privacy features and support access

7.3 International-Transfer Governance

Processor and subprocessor obligations are documented in Data Processing Agreements. Where Chapter V GDPR applies, the relevant transfer mechanism—such as an adequacy decision, Standard Contractual Clauses or Binding Corporate Rules—is documented separately together with any required transfer assessment and supplementary measures. EU data hosting is not treated as conclusive evidence that no international transfer or third-country access occurs.

8. Trust Boundaries

The architecture crosses the customer-to-platform boundary, external-service boundaries and the AWS control-plane boundary. Customer document content, temporary copies, extracted text and AI inference remain inside the regional Veltrixa-controlled AWS data plane. Identity and minimised telemetry cross to separately governed providers, while AWS IAM is assessed as a global control-plane service rather than as region-bound.

Trust Boundaries and External Services




Figure 3. Trust Boundary Diagram. Logical boundaries, external recipients and applicable validation points.

Boundary	Data / Purpose	Primary Controls
Customer ↔ Platform	Document content, upload metadata and retrieval	TLS, authentication, tenant context
Application ↔ AI	Document inference and draft output	Service authentication, ephemeral processing
Application ↔ Storage	Approved documents, metadata and indexes	Tenant-scoped authorisation, encryption
Platform ↔ IdP	Authentication and account events	DPA, federation/MFA, EU configuration
CloudWatch ↔ Datadog EU	Minimised telemetry	Field filtering, TLS, RBAC, retention
Support ↔ Customer Tenant	Exceptional troubleshooting access	Approval, time limit, least privilege, audit log

9. Access Model

Access is governed by tenant-scoped Role-Based Access Control, workload identities and least privilege. Support access is exceptional rather than standing access.

Role	Read	Write	Delete	Administrative Rights
Customer User	Authorised tenant content	Own authorised content	No	No
Customer Administrator	Tenant content	Tenant content and policy	Policy-based	Tenant user and retention administration
Support Engineer	Time-bound, approved content access	No routine write	No	Approved troubleshooting only
Cloud Operations	Infrastructure metadata	Infrastructure configuration	No content deletion	Platform administration
Security Team	Audit and security telemetry	Security cases only	No content deletion	Security monitoring
Automated Service Roles	Task-specific objects and records	Task-specific	Lifecycle-controlled	No interactive login

Role	Read	Write	Delete	Administrative Rights
Recovery Role	Recovery copies under approved event	Restore and reconciliation records	Re-apply validated deletion instructions	Restricted, non-standing recovery administration

10. Third Parties and Transfer Controls

Provider	Service	GDPR Role by Context	Personal Data	Location	Key Controls
AWS	Cloud hosting and backup	Subprocessor for customer content; processor for Veltrixa-controlled data	Documents, metadata, logs, backups	Frankfurt plus service-specific operational access	DPA, regional configuration, KMS, IAM, service review
Datadog	Monitoring	Subprocessor/processor by context	Minimised telemetry; no document content	EU site; support access to validate	DPA, filtering, RBAC, retention, transfer assessment
External Identity Provider	Authentication	Subprocessor/processor by context	Account and authentication data	EU configuration; support access to validate	DPA, federation, MFA, access review

Control	Purpose	Target Status
Data Processing Agreement	Article 28 obligations	Required for each processor/subprocessor relationship
Subprocessor register	Transparency and change management	Required and periodically reviewed
Chapter V mechanism	Legal basis for applicable restricted transfers	Documented where applicable
Transfer assessment	Assess third-country law/practice and supplementary measures	Required where Article 46 is relied upon
Service-specific regional review	Validate storage, processing, support and telemetry locations	Required before production and after material change

11. Privacy Assessment

11.1 Positive Design Characteristics

- Customer document content, temporary copies, extracted text and AI inference workloads are regionally scoped to AWS services in eu-central-1; AWS IAM and other global or control-plane services are assessed separately.
- Document content, metadata, identity data and telemetry are logically separated.
- AI output remains draft until meaningful user review.
- Temporary processing copies are excluded from backup and have fallback expiry.
- External telemetry is minimised and designed to exclude document content.
- Support access is exceptional, approved, time-bound and logged.

11.2 Open Privacy Risks

Risk to Individuals	Likelihood	Impact	Treatment
Unexpected special-category or Article 10 data is processed without appropriate customer governance	Medium	High	Customer guidance, DPA instructions, access restrictions and DPIA screening
Cross-tenant access exposes another customer's documents	Low	High	Tenant isolation, object-level authorisation, testing and audit alerts
Incorrect AI metadata causes misrouting or inappropriate disclosure	Medium	High	Meaningful review, editable output, confidence display and correction workflow

Risk to Individuals	Likelihood	Impact	Treatment
Retention configuration or lifecycle failure causes excessive storage	Medium	Medium	Configuration validation, deletion monitoring and backup-expiry review
Logs or telemetry contain document content or sensitive search terms	Medium	High	Field allowlist, content exclusion, masking and periodic log review
Exceptional support access becomes excessive or insufficiently transparent	Low	High	Approval, time limitation, least privilege and immutable audit trail

11.3 Overall Assessment

The target design includes positive privacy characteristics, including regional processing of customer content, separation of data types, tenant-scoped access, human review and controlled temporary storage. Portfolio-level assurance remains conditional on validation of vendor data flows, support access, log content, backup retention and deletion behaviour. This mapping therefore supports—but does not replace—a DPIA screening, ROPA, vendor review or transfer assessment.

Appendix A – Assumptions and Validation Register

ID	Assumption / Validation Question	Owner	Portfolio Status
VAL-01	Does the selected IdP tenant store identity data and authentication events in the EU, including support access?	Privacy / IAM	Requires evidence
VAL-02	Is Datadog configured on the EU site, and are document bodies, extracted text and sensitive search terms excluded?	Security / Privacy	Requires evidence
VAL-03	Which AWS service features may process service data outside eu-central-1?	Cloud Engineering	Requires service-specific review
VAL-04	Do S3 lifecycle rules remove current and non-current versions within the stated periods?	Cloud Engineering	Requires test evidence
VAL-05	What are the configured RDS automated-backup/PITR period, AWS Backup recovery-point period and controls for any manual or copied snapshots, including failed expiry states?	Cloud Engineering	Requires test evidence
VAL-06	Does the user-review interface provide enough information and authority for meaningful intervention?	Product / Privacy	Requires usability validation
VAL-07	Is privileged support access approved, time-bound and logged in every tenant?	Security / Support	Requires control evidence
VAL-08	Does any job orchestrator, retry queue or dead-letter queue contain document bodies, extracted text or personal identifiers, and what are its retry, visibility and expiry settings?	Cloud Engineering / Privacy	Requires evidence

Appendix B – Glossary

Term	Definition
DFM	Data Flow Mapping
DPIA	Data Protection Impact Assessment
ROPA	Record of Processing Activities
RBAC	Role-Based Access Control
SCC	Standard Contractual Clauses
TIA	Transfer Impact Assessment

Term	Definition
PbD	Privacy by Design
IdP	Identity Provider
SSE-KMS	Server-side encryption using AWS Key Management Service
Article 10 data	Personal data relating to criminal convictions and offences